

Qnn2__extracredit

December 10, 2023

```
[2]: from sklearn.decomposition import PCA
from matplotlib.pyplot import *
from timeit import default_timer as timer

import numpy as np
from nn import NN
from nn import Relu, Linear, SquaredLoss, CELoss
from utils import data_loader, acc, save_plot, loadMNIST, onehot

# Several passes of the training data
def train(model, training_data, dev_data, learning_rate, batch_size, max_epoch):
    X_train, Y_train = training_data['X'], training_data['Y']
    X_dev, Y_dev = dev_data['X'], dev_data['Y']
    for i in range(max_epoch):
        for X,Y in data_loader(X_train, Y_train, batch_size=batch_size,
↪shuffle=True):
            training_loss, grad_Ws, grad_bs = model.compute_gradients(X, Y)
            model.update(grad_Ws, grad_bs, learning_rate)
            dev_acc = acc(model.predict(X_dev), Y_dev)
            print("Epoch {: >3d}/{:}\tloss:{:.5f}\tdev_acc:{:.5f}".
↪format(i+1,max_epoch,training_loss, dev_acc))
        return model

# One pass of the training data
def train_1pass(model, training_data, dev_data, learning_rate, batch_size,
↪print_every=100, plot_every=10):
    X_train, Y_train = training_data['X'], training_data['Y']
    X_dev, Y_dev = dev_data['X'], dev_data['Y']

    num_samples = 0
    print_loss_total = 0
    plot_loss_total = 0

    plot_losses = []
    plot_num_samples = []
    for idx, (X,Y) in enumerate(data_loader(X_train, Y_train,
↪batch_size=batch_size, shuffle=True),1):
```

```

        training_loss, grad_ws, grad_bs = model.compute_gradients(X, Y)
        model.update(grad_ws, grad_bs, learning_rate)
        num_samples += Y.shape[1]
        print_loss_total += training_loss
        plot_loss_total += training_loss

    if idx % print_every == 0:
        dev_acc = acc(model.predict(X_dev), Y_dev)
        print_loss_avg = print_loss_total/print_every
        print_loss_total = 0
        print("#Samples {:>5d}\tloss:{:.5f}\tdev_acc:{:.5f}".
        ↪format(num_samples, print_loss_avg, dev_acc))
    if idx % plot_every == 0:
        plot_loss_avg = plot_loss_total / plot_every
        plot_loss_total = 0
        plot_losses.append(plot_loss_avg)
        plot_num_samples.append(num_samples)

    return model, {"losses":plot_losses, "num_samples":plot_num_samples}

if __name__ == "__main__":
    x_train, label_train = loadMNIST('data/train-images.idx3-ubyte', 'data/
    ↪train-labels.idx1-ubyte')
    x_test, label_test = loadMNIST('data/t10k-images.idx3-ubyte', 'data/
    ↪t10k-labels.idx1-ubyte')
    y_train = onehot(label_train)
    y_test = onehot(label_test)

    dr = PCA()
    #no need to scale data since the features (i.e. pixels) are on the same
    ↪scale
    #need to transpose x since pca requires (N x D)
    full_train = dr.fit_transform(x_train.T)
    #must transform x_test into same space/dimensions as x_train
    #note we cannot perform pca on x_train + x_test together since
    #we are not allowed to use information from x_test during training
    #also we do not perform pca on x_test separated as we would get
    #different principal components (i.e. a projection into a different space)
    #so we use the principal components from x_train
    full_test = dr.transform(x_test.T)

    #get cumulative variance explained
    variances = np.cumsum(dr.explained_variance_ratio_)
    bar(range(1,785), height=variances)
    title("Cumulative Variation Represented by Eigen Vectors")
    xlabel("# eigen vectors")
    ylabel("cumulative % variance")

```

```

show()

#plot projections onto different pairs of principle components
components = [(1,2), (200, 201), (400, 401), (783,784)]
for c1,c2 in components:
    scatter(full_train[:,c1-1], full_train[:,c2-1],s=1)
    title(f"{full_train.shape[0]} training points projected in R2 using
↪PC{c1} and {c2}")
    xlabel(f"Projection onto PC{c1}")
    ylabel(f"Projection onto PC{c2}")
    show()

#what proportion of variance should the reduced data maintain
proportions = [0.10, 0.25, 0.50, 0.80, 0.90, 0.99]
dr_xtrain = []
dr_xtest = []
for p in proportions:
    #num of eigen vectors necessary to explain at least p of the variance
    evs = np.argmax(variances >= p) + 1
    #must take transpose since PCA requires different format of data than NN
    dr_xtrain.append(full_train[:, :evs].T)
    dr_xtest.append(full_test[:, :evs].T)

lr = 1e-2
max_epoch = 20
batch_size = 128

#track train + prediction times
train_times = []
test_times = []
#track test accuracies
accuracies = []
for x_train,x_test in zip(dr_xtrain,dr_xtest):
    #y data can be left untouched, don't need to map into predictor space
    training_data = {"X":x_train, "Y":y_train}
    dev_data = {"X":x_test, "Y":y_test}

    model = NN(Relu(), SquaredLoss(), hidden_layers=[256, 256],
↪input_d=x_train.shape[0], output_d=10)
    model.print_model()

    # model, plot_dict = train_1pass(model, training_data, dev_data, lr,
↪batch_size)
    # save_plot(plot_dict["num_samples"], plot_dict["losses"])
    start_train = timer()
    model = train(model, training_data, dev_data, lr, batch_size, max_epoch)

```

```

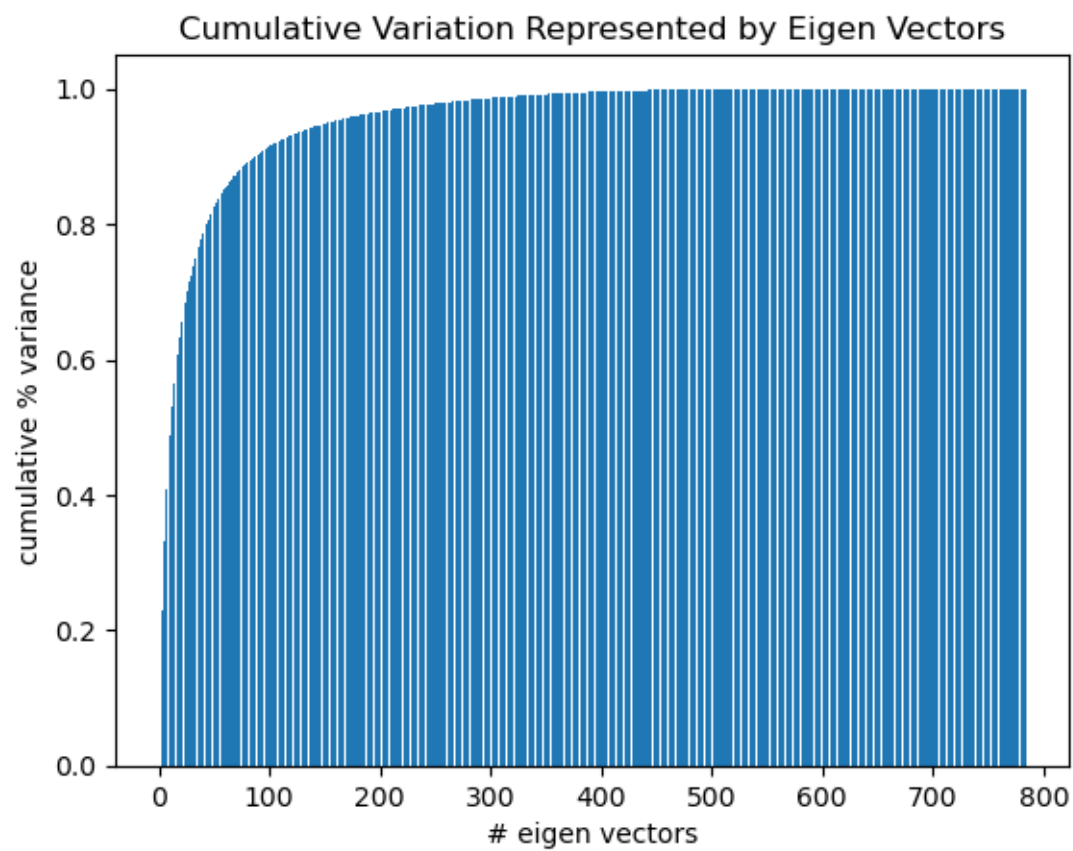
end_train = timer()

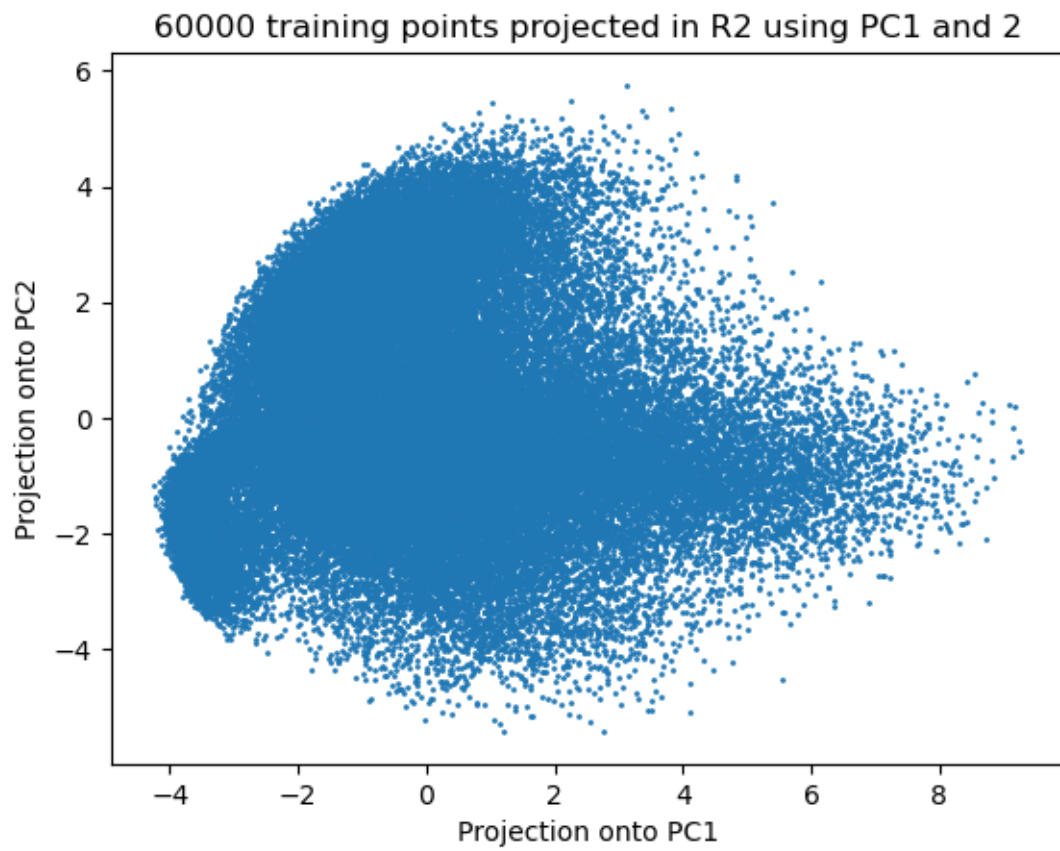
start_test = timer()
accuracy = acc(model.predict(x_test), y_test)
end_test = timer()


train_times.append((end_train-start_train))
test_times.append((end_test-start_test))
accuracies.append(accuracy)


plot(dimensions, train_times, marker='o')
title("NN Train Time vs Dimension")
xlabel("Dimension")
ylabel("Time in Seconds")
show()
plot(dimensions, test_times, marker='o')
title("NN Test Time vs Dimension")
xlabel("Dimension")
ylabel("Time in Seconds")
show()
plot(dimensions, accuracies, marker='o')
title("NN Accuracy on Test Set vs Dimension")
xlabel("Dimension")
ylabel("Accuracy")
show()
plot(proportions, accuracies, marker='o')
title("NN Accuracy on Test Set vs Proportion of Variance")
xlabel("Proportion of Variance")
ylabel("Accuracy")
show()


```











```

activation:Relu
loss function:SquaredLoss
Layer 1 w:(256, 2)      b:(256, 1)
Layer 2 w:(256, 256)    b:(256, 1)
Layer 3 w:(10, 256)     b:(10, 1)
Epoch 1/20  loss:0.36894  dev_acc:0.40800
Epoch 2/20  loss:0.31951  dev_acc:0.41640
Epoch 3/20  loss:0.35647  dev_acc:0.42810
Epoch 4/20  loss:0.36225  dev_acc:0.43190
Epoch 5/20  loss:0.35008  dev_acc:0.43290
Epoch 6/20  loss:0.34070  dev_acc:0.43690
Epoch 7/20  loss:0.29937  dev_acc:0.43570
Epoch 8/20  loss:0.33532  dev_acc:0.43440
Epoch 9/20  loss:0.32132  dev_acc:0.43450
Epoch 10/20 loss:0.32183  dev_acc:0.43870
Epoch 11/20 loss:0.33723  dev_acc:0.43730
Epoch 12/20 loss:0.31511  dev_acc:0.43400
Epoch 13/20 loss:0.36144  dev_acc:0.43850
Epoch 14/20 loss:0.32491  dev_acc:0.43490

```

Epoch	15/20	loss:0.30594	dev_acc:0.43120
Epoch	16/20	loss:0.35773	dev_acc:0.43540
Epoch	17/20	loss:0.31421	dev_acc:0.43490
Epoch	18/20	loss:0.32615	dev_acc:0.43530
Epoch	19/20	loss:0.32410	dev_acc:0.43500
Epoch	20/20	loss:0.32848	dev_acc:0.43560

activation:Relu

loss function:SquaredLoss

Layer 1 w:(256, 4) b:(256, 1)

Layer 2 w:(256, 256) b:(256, 1)

Layer 3 w:(10, 256) b:(10, 1)

Epoch	1/20	loss:0.33211	dev_acc:0.56230
Epoch	2/20	loss:0.26983	dev_acc:0.58580
Epoch	3/20	loss:0.29336	dev_acc:0.60050
Epoch	4/20	loss:0.27551	dev_acc:0.60600
Epoch	5/20	loss:0.25036	dev_acc:0.60860
Epoch	6/20	loss:0.24512	dev_acc:0.61140
Epoch	7/20	loss:0.22633	dev_acc:0.61210
Epoch	8/20	loss:0.25574	dev_acc:0.61400
Epoch	9/20	loss:0.25584	dev_acc:0.61360
Epoch	10/20	loss:0.24913	dev_acc:0.61730
Epoch	11/20	loss:0.25191	dev_acc:0.62020
Epoch	12/20	loss:0.26015	dev_acc:0.61810
Epoch	13/20	loss:0.27953	dev_acc:0.62270
Epoch	14/20	loss:0.26127	dev_acc:0.62270
Epoch	15/20	loss:0.23365	dev_acc:0.62570
Epoch	16/20	loss:0.23718	dev_acc:0.62480
Epoch	17/20	loss:0.20384	dev_acc:0.62640
Epoch	18/20	loss:0.23292	dev_acc:0.62990
Epoch	19/20	loss:0.25073	dev_acc:0.63280
Epoch	20/20	loss:0.25028	dev_acc:0.63040

activation:Relu

loss function:SquaredLoss

Layer 1 w:(256, 11) b:(256, 1)

Layer 2 w:(256, 256) b:(256, 1)

Layer 3 w:(10, 256) b:(10, 1)

Epoch	1/20	loss:0.21811	dev_acc:0.75630
Epoch	2/20	loss:0.19837	dev_acc:0.81490
Epoch	3/20	loss:0.15654	dev_acc:0.83790
Epoch	4/20	loss:0.17383	dev_acc:0.84930
Epoch	5/20	loss:0.10895	dev_acc:0.85700
Epoch	6/20	loss:0.11517	dev_acc:0.86380
Epoch	7/20	loss:0.13261	dev_acc:0.86950
Epoch	8/20	loss:0.10374	dev_acc:0.87520
Epoch	9/20	loss:0.14593	dev_acc:0.87860
Epoch	10/20	loss:0.10694	dev_acc:0.88150
Epoch	11/20	loss:0.12225	dev_acc:0.88500
Epoch	12/20	loss:0.10560	dev_acc:0.88670

Epoch	13/20	loss:0.10249	dev_acc:0.88920
Epoch	14/20	loss:0.08595	dev_acc:0.89100
Epoch	15/20	loss:0.12218	dev_acc:0.89160
Epoch	16/20	loss:0.12055	dev_acc:0.89310
Epoch	17/20	loss:0.06458	dev_acc:0.89480
Epoch	18/20	loss:0.13032	dev_acc:0.89680
Epoch	19/20	loss:0.10664	dev_acc:0.89790
Epoch	20/20	loss:0.10433	dev_acc:0.89820

activation:Relu

loss function:SquaredLoss

Layer 1 w:(256, 44) b:(256, 1)

Layer 2 w:(256, 256) b:(256, 1)

Layer 3 w:(10, 256) b:(10, 1)

Epoch	1/20	loss:0.32145	dev_acc:0.81350
Epoch	2/20	loss:0.21285	dev_acc:0.87480
Epoch	3/20	loss:0.17348	dev_acc:0.89660
Epoch	4/20	loss:0.16596	dev_acc:0.90950
Epoch	5/20	loss:0.13465	dev_acc:0.91640
Epoch	6/20	loss:0.12378	dev_acc:0.92290
Epoch	7/20	loss:0.13350	dev_acc:0.92530
Epoch	8/20	loss:0.09641	dev_acc:0.92960
Epoch	9/20	loss:0.11166	dev_acc:0.93210
Epoch	10/20	loss:0.11379	dev_acc:0.93330
Epoch	11/20	loss:0.09142	dev_acc:0.93640
Epoch	12/20	loss:0.11195	dev_acc:0.93690
Epoch	13/20	loss:0.08888	dev_acc:0.93760
Epoch	14/20	loss:0.09207	dev_acc:0.94000
Epoch	15/20	loss:0.09940	dev_acc:0.93940
Epoch	16/20	loss:0.08760	dev_acc:0.94200
Epoch	17/20	loss:0.10389	dev_acc:0.94160
Epoch	18/20	loss:0.10943	dev_acc:0.94270
Epoch	19/20	loss:0.08941	dev_acc:0.94380
Epoch	20/20	loss:0.09978	dev_acc:0.94480

activation:Relu

loss function:SquaredLoss

Layer 1 w:(256, 87) b:(256, 1)

Layer 2 w:(256, 256) b:(256, 1)

Layer 3 w:(10, 256) b:(10, 1)

Epoch	1/20	loss:0.32317	dev_acc:0.78410
Epoch	2/20	loss:0.25379	dev_acc:0.85300
Epoch	3/20	loss:0.19423	dev_acc:0.87890
Epoch	4/20	loss:0.14025	dev_acc:0.89390
Epoch	5/20	loss:0.17181	dev_acc:0.90150
Epoch	6/20	loss:0.15374	dev_acc:0.90650
Epoch	7/20	loss:0.13893	dev_acc:0.91110
Epoch	8/20	loss:0.11867	dev_acc:0.91460
Epoch	9/20	loss:0.13896	dev_acc:0.91950
Epoch	10/20	loss:0.10716	dev_acc:0.92220

Epoch	11/20	loss:0.15091	dev_acc:0.92370
Epoch	12/20	loss:0.12702	dev_acc:0.92620
Epoch	13/20	loss:0.10267	dev_acc:0.92710
Epoch	14/20	loss:0.11721	dev_acc:0.92960
Epoch	15/20	loss:0.08857	dev_acc:0.93040
Epoch	16/20	loss:0.09396	dev_acc:0.93190
Epoch	17/20	loss:0.09741	dev_acc:0.93360
Epoch	18/20	loss:0.09721	dev_acc:0.93530
Epoch	19/20	loss:0.09332	dev_acc:0.93530
Epoch	20/20	loss:0.10726	dev_acc:0.93650

activation:Relu

loss function:SquaredLoss

Layer 1 w:(256, 331) b:(256, 1)

Layer 2 w:(256, 256) b:(256, 1)

Layer 3 w:(10, 256) b:(10, 1)

Epoch	1/20	loss:0.29699	dev_acc:0.78120
Epoch	2/20	loss:0.23514	dev_acc:0.84640
Epoch	3/20	loss:0.20420	dev_acc:0.87660
Epoch	4/20	loss:0.16098	dev_acc:0.89210
Epoch	5/20	loss:0.14342	dev_acc:0.90270
Epoch	6/20	loss:0.13605	dev_acc:0.91130
Epoch	7/20	loss:0.17145	dev_acc:0.91540
Epoch	8/20	loss:0.11264	dev_acc:0.92070
Epoch	9/20	loss:0.14375	dev_acc:0.92280
Epoch	10/20	loss:0.11768	dev_acc:0.92550
Epoch	11/20	loss:0.14149	dev_acc:0.92810
Epoch	12/20	loss:0.11686	dev_acc:0.92940
Epoch	13/20	loss:0.10794	dev_acc:0.93140
Epoch	14/20	loss:0.10906	dev_acc:0.93240
Epoch	15/20	loss:0.09770	dev_acc:0.93410
Epoch	16/20	loss:0.10696	dev_acc:0.93580
Epoch	17/20	loss:0.09318	dev_acc:0.93630
Epoch	18/20	loss:0.08980	dev_acc:0.93710
Epoch	19/20	loss:0.12739	dev_acc:0.93790
Epoch	20/20	loss:0.09715	dev_acc:0.93930

